



# Exercise - 8

05.04.2021

**CED17I017**

**FIROZ MOHAMMAD**


## Write on how a SNULL(Simple Network Utility for Loading Localities) works, need not execute, just soft copy is sufficient.

SNULL – Simple Network Utility for Loading Localities,

- driver of the network device
- driver that does not talk to the "actual" devices
- works like a loopback device

### How snull driver work ?

- The *snull* module creates two interfaces.
- **snullnet0** - the network associated with the **sn0** interface,
- **snullnet1** - the network associated with the **sn1** interface,
- **snullnet0** is the network that is connected to the **sn0** interface. Similarly, **snullnet1** is the network connected to **sn1**. The addresses of these networks should differ only in the least significant bit of the third octet. These networks must have 24-bit netmasks.
- **local0** is the IP address assigned to the **sn0** interface; it belongs to **snullnet0**. The address associated with **sn1** is **local1**. **local0** and **local1** must differ in the least significant bit of their third octet and in the fourth octet.
- **remote0** is a host in **snullnet0**, and its fourth octet is the same as that of **local1**. Any packet sent to **remote0** reaches **local1** after its network address has been modified by the interface code. The host **remote1** belongs to **snullnet1**, and its fourth octet is the same as that of **local0**.



Here are possible values for the network numbers. Once you put these lines in */etc/networks*, you can call your networks by name. The values were chosen from the range of numbers reserved for private use.

**snullnet0**    192.168.0.0

**snullnet1**    192.168.1.0

the following are possible host numbers to put into */etc/hosts*:

**192.168.0.1**   *local0*

**192.168.0.2**   *remote0*

**192.168.1.2**   *local1*

**192.168.1.1**   *remote1*

Now , we can ping :

```
user: ~$ ping -c 2 remote0
```

```
64 bytes from 192.168.0.99: icmp_seq=0 ttl=64 time=1.6 ms
```

```
64 bytes from 192.168.0.99: icmp_seq=1 ttl=64 time=0.9 ms
```

```
2 packets transmitted, 2 packets received, 0% packet loss
```

```
user: ~$ ping -c 2 remote1
```

```
64 bytes from 192.168.1.88: icmp_seq=0 ttl=64 time=1.8 ms
```

```
64 bytes from 192.168.1.88: icmp_seq=1 ttl=64 time=0.9 ms
```

```
2 packets transmitted, 2 packets received, 0% packet loss
```

## Let's deep dive to SNULL module working :

### Entry and Exit function :

The Linux network device driver basically makes a fuss about the `net_device` structure. The entry function of the program is mainly to allocate and register the `net_device` structure. Correspondingly, the exit function needs to cancel the logout and release the `net_device` structure. This experiment will eventually use the ping command to test the sending and receiving of two network devices, so first define two pointers to `net_device`:

```
struct net_device *snnull_devs[2];
```

### Entry Function :

```
static __init int snnull_init_module(void)
```

```

{
    int result, i, ret = -ENOMEM;

    /* Allocate the devices */
    snull_devs[0] = alloc_netdev(sizeof(struct snull_priv), "sn%d",
snull_init);
    snull_devs[1] = alloc_netdev(sizeof(struct snull_priv), "sn%d",
snull_init);
    if (snull_devs[0] == NULL || snull_devs[1] == NULL)
        goto out;

    ret = -ENODEV;
    for (i = 0; i < 2; i++)
        if ((result = register_netdev(snull_devs[i])))
            printk("snull: error %i registering device \"%s\"\n",
                result, snull_devs[i]->name);
        else
            ret = 0;
out:
    if (ret)
        snull_cleanup();
    return ret;
}
module_init(snull_init_module);

```

Alloc\_netdev() function is used to allocate net\_device structure:

```

struct net_device *alloc_netdev(int sizeof_priv, const char
*name,
                                void (*setup)(struct
net_device *));

```

The first parameter of the function is the size of the private data, the second parameter is the device name, and the third parameter setup points to a function that takes struct net\_device\* as a parameter and is used to initialize some members of net\_device.

The private data of the program defines snull\_priv (the usage of spin lock members is not repeated):

```
struct snull_priv {
    struct net_device_stats stats;
    int status;
    struct snull_packet *ppool;
    struct snull_packet *rx_queue;
    int rx_int_enabled;
    int tx_packetlen;
    u8 *tx_packetdata;
    struct sk_buff *skb;
    spinlock_t lock;
};
```

The snull\_packet structure is used to save the sent and received data:

```
struct snull_packet {
    struct snull_packet *next;
    struct net_device *dev;
    int datalen;
    u8 data[ETH_DATA_LEN];
};
```

The series of operation functions defined for struct snull\_packet \*ppool are as follows:

```
int pool_size = 8;
module_param(pool_size, int, 0);

/*
 * Allocate 8 snull_packets and add them to a linked list
 */
void snull_setup_pool(struct net_device *dev)
{
```

```

    struct snull_priv *priv = netdev_priv(dev);
    int i;
    struct snull_packet *pkt;

    priv->ppool = NULL;
    for (i = 0; i < pool_size; i++) {
        pkt = kmalloc (sizeof (struct snull_packet), GFP_KERNEL);
        if (pkt == NULL) {
            printk (KERN_NOTICE "Ran out of memory allocating packet
pool\n");
            return;
        }
        pkt->dev = dev;
        pkt->next = priv->ppool;
        priv->ppool = pkt;
    }
}

/*
 * Release the linked list composed of snull_packet
 */
void snull_teardown_pool(struct net_device *dev)
{
    struct snull_priv *priv = netdev_priv(dev);
    struct snull_packet *pkt;

    while ((pkt = priv->ppool)) {
        priv->ppool = pkt->next;
        kfree (pkt);
    }
}

/*
 * Take a snull packet from the linked list for sending
 */
struct snull_packet *snull_get_tx_buffer(struct net_device *dev)
{
    struct snull_priv *priv = netdev_priv(dev);
    unsigned long flags;
    struct snull_packet *pkt;

    spin_lock_irqsave(&priv->lock, flags);

```

```

    pkt = priv->ppool;
    priv->ppool = pkt->next;
    if (priv->ppool == NULL) {
        printk (KERN_INFO "Pool empty\n");
        netif_stop_queue(dev);
    }
    spin_unlock_irqrestore(&priv->lock, flags);
    return pkt;
}

/*
 * Put a snull packet back to the linked list
 */
void snull_release_buffer(struct snull_packet *pkt)
{
    unsigned long flags;
    struct snull_priv *priv = netdev_priv(pkt->dev);

    spin_lock_irqsave(&priv->lock, flags);
    pkt->next = priv->ppool;
    priv->ppool = pkt;
    spin_unlock_irqrestore(&priv->lock, flags);
    if (netif_queue_stopped(pkt->dev) && pkt->next == NULL)
        netif_wake_queue(pkt->dev);
}

```

rx\_queue is the receiving queue of the device, the operation functions are as follows:

```

void snull_enqueue_buf(struct net_device *dev, struct snull_packet *pkt)
{
    unsigned long flags;
    struct snull_priv *priv = netdev_priv(dev);

    spin_lock_irqsave(&priv->lock, flags);
    pkt->next = priv->rx_queue; /* FIXME - misorders packets */
    priv->rx_queue = pkt;
    spin_unlock_irqrestore(&priv->lock, flags);
}

struct snull_packet *snull_dequeue_buf(struct net_device *dev)

```



```

{
    struct snull_priv *priv = netdev_priv(dev);
    struct snull_packet *pkt;
    unsigned long flags;

    spin_lock_irqsave(&priv->lock, flags);
    pkt = priv->rx_queue;
    if (pkt != NULL)
        priv->rx_queue = pkt->next;
    spin_unlock_irqrestore(&priv->lock, flags);
    return pkt;
}

```

rx\_int\_enabled is used to enable or disable the receiving interrupt (the interrupt is not a real hardware interrupt, but is simulated by software), the operation function is as follows:

```

/*
 * Enable and disable receive interrupts.
 */
static void snull_rx_ints(struct net_device *dev, int enable)
{
    struct snull_priv *priv = netdev_priv(dev);
    priv->rx_int_enabled = enable;
}

```

The setup parameter of alloc\_netdevice() points to a function that is used to initialize other members of net\_device. The snull\_init function is used in the program:

```

static int timeout = SNULL_TIMEOUT; //SNULL_TIMEOUT is defined as 5 in
snnull.h
module_param(timeout,int,0);

static const struct net_device_ops snull_dev_ops = {
    .ndo_open = snull_open,
    .ndo_stop = snull_release,
    .ndo_start_xmit = snull_tx,
    .ndo_do_ioctl = snull_ioctl,
}

```

```

        .ndo_get_stats = snull_stats,
        .ndo_tx_timeout = snull_tx_timeout,
    };

    static const struct header_ops snull_header_ops= {
        .create      = snull_header,
        .rebuild     = snull_rebuild_header,
        .cache       = NULL,
    };

    void snull_init(struct net_device *dev)
    {
        struct snull_priv *priv = NULL;

        ether_setup(dev);

        dev->netdev_ops = &snull_dev_ops;
        dev->header_ops = &snull_header_ops;
        dev->watchdog_timeo = timeout;
        dev->flags |= IFF_NOARP; //Forbid ARP
        dev->features |= NETIF_F_NO_CSUM;

        priv = netdev_priv(dev);
        memset(priv, 0, sizeof(struct snull_priv));
        spin_lock_init(&((struct snull_priv *)priv)->lock);
        snull_rx_ints(dev, 1);          /* enable receive interrupts */
        snull_setup_pool(dev);

        return;
    }

```

snull\_init() first initializes the members in the net\_device structure: call the ether\_setup() function to assign default values to many members in the net\_device; Next, set the device method, because the kernel version is different, the device method in the book is directly included in the net\_device structure, which is required here Modified, from the code, you can see that the operations such as open and release are encapsulated in the netdev\_ops, header\_ops structure, and the specific implementation of the device methods such as open, send, and receive

will be further explained later. The `watchdog_timeo` member sets the transmission timeout period, followed by some flags, indicating that ARP is forbidden and force packet verification.

After setting up the `net_device` member, use `netdev_priv()` to obtain the private data of the device and set the spin lock, receive enable and data buffer. `netdev_priv()` is dedicated to access private data of network devices:


```
void *netdev_priv(struct net_device *dev);
```

After allocating and initializing the `net_device` structure, we also need to register this device with `register_netdev`, and register the function prototype for unregistering the network device:

```
int register_netdev(struct net_device *dev);  
void unregister_netdev(struct net_device *dev);
```

## Exit Function :

```
static __exit void snull_cleanup(void)  
{  
    int i;  
  
    for (i=0; i<2;i++)  
    {  
        if(snull_devs[i])  
        {  
            unregister_netdev(snull_devs[i]);  
            snull_tear_down_pool(snull_devs[i]);  
            free_netdev(snull_devs[i]);  
        }  
    }  
    return;  
}  
module_exit(snull_cleanup);
```



The operation of the exit function is the opposite of the entry, mainly to unregister the device, release the data buffer and release the memory occupied by the device.